

Performance optimization sins

Real-life examples



Aliaksandr Valialkin
VictoriaMetrics founder and core developer

Hello

- I'm Aliaksandr Valialkin (aka @valyala)
- The author of fasthttp, fastjson, quicktemplate, etc.
- I'm fond of performance optimizations
- [VictoriaMetrics](#) founder and core developer

When do we need performance optimizations?

- Slow programs
- High costs
- Benchmark games

Performance optimization sins

- Better performance isn't free
- The price may be too high

Performance optimization sins

- Fasthttp
- Fastjson
- VictoriaMetrics
- Standard Go packages
- CPUs

<https://github.com/valyala/fasthttp>

Fasthttp

- Response buffering
- Deferring HTTP headers' parsing
- Slices instead of maps for (key -> value) entries
- RequestCtx re-use
- DNS caching

Fasthttp response buffering

- Standard HTTP protocol flow

```
func ProcessHTTPConn(c net.Conn) {  
    for {  
        req, closed := readRequest(c)  
        if closed {  
            return  
        }  
        resp := processRequest(req)  
        writeResponse(c, resp)  
    }  
}
```


Fasthttp response buffering

- Protocol flow optimized for HTTP pipelining

```
func ProcessPipelinedHTTPConn(c net.Conn) {  
    br := bufio.NewReader(c)  
    bw := bufio.NewWriter(c)  
    for {  
        req, closed := readRequest(br)  
        if closed {  
            bw.Flush()  
            return  
        }  
        resp := processRequest(req)  
        writeResponse(bw, resp)  
        if br.Buffered() == 0 {  
            // Send buffered responses to client, since no more buffered requests  
            bw.Flush()  
        }  
    }  
}
```

Fasthttp response buffering benefits

- Reduces the number of **recv** syscalls, since multiple small requests may be read from conn via a single **recv** syscall
- Reduces the number of **send** syscalls, since multiple responses may be flushed to conn via a single **send** syscall

Fasthttp response buffering sins

- Delays responses for slow pipelined requests
- Doesn't provide any speedup for real-world HTTP servers, since modern web-browsers don't use HTTP pipelining due to historical bugs. See https://en.wikipedia.org/wiki/HTTP_pipelining#Implementation_status
- Provides 2x speedup for Techempower plaintext benchmark. See <https://github.com/TechEmpower/FrameworkBenchmarks/issues/4410>

Deferring HTTP headers' parsing

- Fasthttp defers parsing of HTTP headers until they are needed
- It just searches for the “end of headers” marker in the input buffer and puts unparsed headers into a byte slice
- The byte slice is parsed into HTTP headers structure on the first access

Deferring HTTP headers' parsing

- The following code skips parsing HTTP headers:

```
fasthttp.ListenAndServe(":8080", func(ctx *fasthttp.RequestCtx) {  
    ctx.WriteString("Hello, world!")  
})
```

Deferring HTTP headers' parsing

The following code parses HTTP headers:

```
fasthttp.ListenAndServe(":8080", func(ctx *fasthttp.RequestCtx) {  
    clientAddr := ctx.Request.Header.Peek("X-Forwarded-For")  
    fmt.Fprintf(ctx, "client addr: %q", clientAddr)  
})
```

Deferring HTTP headers' parsing benefits

- If request handler doesn't access HTTP headers, then CPU time is saved

Deferring HTTP headers' parsing sins

- Real-world HTTP handlers usually consult HTTP headers
- Techempower plaintext benchmark is the only user that benefits from this optimization :)

Slices instead of maps for (key -> value)

- Fasthttp uses slices instead of maps for storing the following entities:
 - HTTP headers
 - Query args
 - Cookies

Slices instead of maps for (key -> value)

- (key->value) slice structs:

```
type kv struct {  
    key    []byte  
    value []byte  
}
```

```
type sliceMap []kv
```

Slices instead of maps for (key -> value)

- How to add an entry to **sliceMap** and re-use memory

```
func (sm *sliceMap) Add(k, v []byte) {  
    kvs := *sm  
    if cap(kvs) > len(kvs) {  
        kvs = kvs[:len(kvs)+1]  
    } else {  
        kvs = append(kvs, kv{})  
    }  
    kv := &kvs[len(kvs)-1]  
    kv.key = append(kv.key[:0], k...)  
    kv.value = append(kv.value[:0], v...)  
    *sm = kvs  
}
```

Slices instead of maps for (key -> value)

- How to get value for the given key from **sliceMap**

```
func (sm sliceMap) Get(k string) []byte {  
    for i := range sm {  
        kv := &sm[i]  
        if string(kv.key) == k {  
            return kv.value  
        }  
    }  
    return nil  
}
```

Slices instead of maps for (key -> value): benefits

- Usually the number of entries in headers, query args or cookies is quite small (less than 10)
- Slices provide better performance comparing to maps for this case
- Slices allow memory re-use
- Slices save the original order of added items

Slices instead of maps for (key->value): sins

- **sliceMap.Get** has **$O(N)$** complexity, while standard map has **$O(1)$** complexity
- **sliceMap** is vulnerable to memory fragmentation on re-use:
 - Write **hugeValue** into **sliceMap**
 - Now **sliceMap** occupies at least **len(hugeValue)** bytes of memory when re-used
- The value returned from **sliceMap.Get** is valid only until the **sliceMap** is re-used

RequestCtx re-use

```
func processConn(c net.Conn, requestHandler RequestHandler) {  
    // The ctx is re-used across requestHandler calls  
    ctx := AcquireRequestCtx()  
    defer ReleaseRequestCtx(ctx)  
    for {  
        ctx.Reset()  
        readRequest(c, ctx)  
        requestHandler(ctx)  
        writeResponse(c, ctx)  
    }  
}
```

RequestCtx re-use: benefits

- Reduced memory allocations, since the **RequestCtx** memory is re-used
- Better performance, since **RequestCtx** is already in CPU cache on re-use

RequestCtx re-use: sins

- Easy to shoot in the foot by holding references to **RequestCtx** contents after returning from **RequestHandler**
- Possible memory fragmentation/calification on **RequestCtx** re-use

DNS caching

- Fasthttp caches (**host -> IP**) entries for a minute

```
func resolveHost(host string) net.IP {  
    e := dnsCache[host]  
    if e != nil && time.Since(e.resolveTime) < time.Minute {  
        // Fast path - return the ip from cache.  
        return e.ip  
    }  
    // Slow path - really resolve the host to ip and put it to cache  
    ip := reallyResolveHost(host)  
    dnsCache[host] = ip  
    return ip  
}
```

DNS caching benefits

- Reduced load on DNS subsystem
- Faster dials to remote hosts

DNS caching sins

- Breaks on frequent DNS changes (when Kubernetes restarts pods)

<https://github.com/valyala/fastjson>

Fastjson

- Memory re-use
- Fast string unescaping
- Custom parser for integers and floats

Fastjson memory re-use

- **fastjson.Parser** owns all the JSON data structure
- **fastjson.Parser** re-uses the memory for JSON data structure

```
var p fastjson.Parser
v, err := p.Parse(`{"foo": "bar"}`) // v belongs to p
...
b := v.GetStringBytes("foo") // b also belongs to p
...
vv, err := p.Parse(`[1,2,3]`)
// vv overwrites v contents, so v and b become invalid
```

Fastjson memory re-use: benefits

- Reduced memory allocations
- Improved performance, since the memory remains in CPU caches across **Parse** invocations

Fastjson memory re-use: sins

- High memory usage after parsing random big JSON objects
- “Shoot in the foot” API - all the JSON structures returned by **Parser** cannot be used after the next **Parse** call. Recursively

Fast string unescaping

```
func unescapeJSONString(s string) string {  
    n := strings.IndexByte(s, '\\')  
    if n < 0 {  
        // Fast path - the string has no escape chars.  
        return s  
    }  
    // Slow path - unescape every char in s  
    return unescapeJSONStringSlow(s)  
}
```

Fast string unescaping: benefits

- Works fast for strings without escape chars

Fast string unescaping: sins

- Works slow on strings with escaped chars
- Uneven performance on mixed strings

Custom parser for integer and floats

- Fastjson doesn't use **strconv.ParseFloat** and **strconv.ParseInt**
- It uses custom functions optimized for speed

Custom parser for integer and floats: benefits

- Faster than the corresponding functions from **strconv**

Custom parser for integer and floats: sins

- It falls back to **strconv** for too big numbers -> slower performance
- It returns 0 for invalid numbers
- It may work improperly for corner cases

<https://victoriametrics.com>

VictoriaMetrics

- Hand-written protobuf parsing with memory re-use

Hand-written protobuf parsing

- VictoriaMetrics accepts protobuf via [Prometheus remote write API](#)
- Initially we used parser generated by standard protobuf generator
- It was allocating like a hell, so it had been rewritten to zero-alloc mode
- Now protobuf parsers re-use the provided structs

Hand-written protobuf parser

```
type TimeSeries struct {  
    Labels []Label  
    Samples []Sample  
}
```

```
type Label struct {  
    Name []byte  
    Value []byte  
}
```

```
type Sample struct {  
    Value float64  
    Timestamp int64  
}
```

Hand-written protobuf parser

- Diff examples:

- Name string

- + Name []byte

- m.Name = string(dAtA[iNdEx:postIndex])

- + m.Name = dAtA[iNdEx:postIndex]

Hand-written protobuf parser

- Diff examples:

```
- Labels  []*Label
+ Labels  []Label
```

```
- m.Labels = append(m.Labels, &Label{})
+ if cap(m.Labels) > len(m.Labels) {
+     m.Labels = m.Labels[:len(m.Labels)+1]
+ } else {
+     m.Labels = append(m.Labels, Label{})
+ }
```

Hand-written protobuf parsing: benefits

- Zero allocations
- Better performance

Hand-written protobuf parsing: sins

- The hand-written code is fragile
- It is hard to update the code if Prometheus remote write API changes
- The parsed struct cannot be referenced after the next Unmarshal call

Standard Go packages

Standard Go packages

- Pools for small numbers
- **sync.Pool** memory fragmentation

Pools for small numbers

- Go doesn't allocate on `strconv.Itoa(smallNum)`:

```
func FormatInt(i int64, base int) string {  
    if 0 <= i && i < nSmalls && base == 10 {  
        return small(int(i))  
    }  
    _, s := formatBits(nil, uint64(i), base, i < 0, false)  
    return s  
}
```

```
func small(i int) string {  
    if i < 10 {  
        return digits[i : i+1]  
    }  
    return smallsString[i*2 : i*2+2]  
}
```

Pools for small numbers: benefits

- Faster performance for small numbers (up to 99)
- Zero memory allocations for small numbers

Pools for small numbers: sins

- Uneven performance for mixed numbers
- Slower performance for bigger numbers

sync.Pool memory fragmentation

- sync.Pool allows re-using objects with the underlying memory
- Benefit: reduced memory allocations and improved performance
- Sin: possible high memory usage:
 - Suppose **sync.Pool** contains small byte slices
 - Put a big byte slice into **sync.Pool**
 - Now the big byte slice wastes memory when obtained from the pool, since only a small amount of allocated memory is really used

CPU's

CPUs

- Modern CPUs execute instructions in multi-stage pipeline
- This improves performance
- The pipeline is reset on each conditional branch, leading to delay
- CPU makers added prediction block and speculative execution, which may go beyond branches
- Benefit: CPU runs full speed until the first branch mispredict
- Sins: **Meltdown** and **Spectre**-like vulnerabilities

Conclusion

Conclusion

- Performance optimization is full of hard decisions:
 - Speed vs clarity
 - Speed vs simplicity
 - Speed vs nice API
 - Speed vs consistent performance for all the cases
 - Speed vs consistent memory usage
 - Speed vs precision
 - Speed vs lower vulnerability risk
- There is no silver bullet for making perfect decision
- Choose wisely

Questions?



Aliaksandr Valialkin
VictoriaMetrics founder and core developer